

General Research Notes

Zachary Ross

June 16, 2022

Contents

I Dated Findings	3
February 2, 2022	4
February 9, 2022	4
February 23, 2022	4
II Papers (Deprecated)	6
1 Algorithms and Theory	8
1.1 Monte Carlo and Las Vegas Randomized Algorithms for Systems and Control*: An Introduction [25]	8
1.2 Approximation Algorithms for k-Anonymity [1]	8
1.3 Improved Approximation Algorithms for Maximum Cut and Satis- fiability Problems Using Semidefinite Programming [8]	9
1.3.1 Analysis	10
1.4 Treewidth, partial k -trees, and chordal graphs [10]	11
2 Deep Learning	13
2.1 Efficient Per-Example Gradient Computations [9]	13
2.1.1 Commentary	14
2.2 Understanding gradient clipping in incremental gradient methods [21]	14
2.3 Pruning neural networks without any data by iteratively conserving synaptic flow [24]	17
2.3.1 Pruning Algorithms	17
2.3.2 Synaptic Saliency	17
2.3.3 Algorithm	18
2.3.4 Commentary	19
2.4 Learning Low-rank Deep Neural Networks via Singular Vector Or- thogonality Regularization and Singular Value Sparsification [28] .	19
2.5 Scalable Adaptive Stochastic Optimization Using Random Projec- tions [15]	21
2.6 A Mini-Block Natural Gradient Method for Deep Neural Networks [3]	23
2.6.1 Preconditioning	24
3 Finance	25
3.1 An Introduction to Quantitative Finance [4]	25
3.1.1 Interest Rates	25

CONTENTS	2
3.1.2 Forward Contracts and Forward Prices	28
3.2 A Random Walk Down Wall Street [18]	29
III Subjects	30
4 Probability and Statistics	31
5 Stochastic Optimization	32
5.1 Quasi-Newton Methods	32
5.1.1 BFGS	32
6 Federated Learning	35
6.1 Update Methods	35
6.1.1 Structured Updates	36
6.1.2 Sketched Updates	36
7 Differential Privacy	38
7.1 (ϵ, δ) -Differential Privacy	38
7.1.1 Mechanisms	38
7.2 Rényi Differential Privacy	39
8 Jacobian Orthogonalization in Neural Networks	41
8.1 Linear Layers	41
8.1.1 Householder Reflectors	41
8.1.2 Orthogonality Regularization	42
8.1.3 Cayley Transform	42
8.2 Convolutional Layers	43
8.2.1 Explicit Convolution Construction	43
8.2.2 Circularly Padded Convolutions	43
9 Random Projections	45

Part I

Dated Findings

February 2, 2022

In the work presented to me, Dr. Lee primarily uses pruning for gradient masking rather than parameter masking, although most studies I've read seem to have done the opposite. I initially assumed that given the standard linear layer function with masking

$$\mathbf{z} = (\mathbf{M} \odot \mathbf{W})\mathbf{h} + \mathbf{b} \quad (1)$$

the effect of the mask would still result in non-zero values in the gradient of the matrix. This is falsified by the derivation

$$\frac{\partial L}{\partial W_{i,j}} = \frac{\partial L}{\partial z_i} \frac{\partial z_i}{\partial W_{ij}} = M_{ij} \frac{\partial L}{\partial z_i} h_j \quad (2)$$

or in vectorized form we have that

$$\frac{\partial L}{\partial \mathbf{W}} = \mathbf{M} \odot \left(\frac{\partial L}{\partial \mathbf{z}} (\mathbf{h})^\top \right) \quad (3)$$

which implies that the use of a forward mask guarantees that of a backward. The method used instead in Dr. Lee's examples was $\mathbf{z} = \mathbf{W}\mathbf{h} + \mathbf{b}$ with masking only applied during the gradient (i.e. that of Equation 3). So what needs to be tested is whether or not we can apply masking solely during the forward phase i.e. implement Equation 1 with builds its gradient without the mask.

February 9, 2022

Thorough testing has disproved our previous hypothesis. Strangely enough, the gradient masking is much more effective than weight masking for some odd reason.

February 23, 2022

In the following discussion, I use beats per minute (BPM) to refer to heart rate. I only do it this way because I've already written it out and don't feel like changing it.

I've been pondering a bit recently on a hypothesis I've had, that being the perception of time is relative to the individual based on several factors. At first I believed that it's connected to BPM, but this was disproved by my addiction to coffee. I've noticed this relativity even in my day to day, some days seem faster than others while certain activities seem to be slower. For example, under the influence of coffee I feel that my day goes by much faster, which leads me to believe the effects of caffeine are misunderstood (which I will go into later). After or during an intense workout, my perception of time seems to slow way the hell down. I notice this because I'll be listening to music, and maybe it has to do with my percussionist background with intense focus on steady timing, but the songs pass by much slower. This is incredibly curious.

On other days where I don't quite feel myself, I notice instead that the music tends to be much quicker, often times leading me to feel like I'm missing a lot.

For the coffee thing, maybe its affect on BPM does slow down perception of time, but caffeine could effectively and periodically cut out clips of time, making

us believe that time is moving faster. I believe this affects productivity because it allows us to move from task to task seemingly "quicker". I don't believe it fully alters our productivity; rather, it seems like it allows us to reach the reward much quicker and thereby reinforce the thought process that we are "getting shit done". This activity/reward behavior most likely becomes a self-fulfilling cycle as the momentum continues, motivating us to continue doing tasks.

I've also considered how we would test these hypotheses because it's rather elusive. Say we tried to get a set of people to listen to a metronome prior to an energy intensive activity and post said activity and asked them to compare the two for which one is faster. For one, how would we get reliable information? If we were just to pick random people from a population, they may not even be capable of discerning between two tempos. I suppose we could do a preliminary exam where we vet out those who could not tell the difference even between minor changes in tempo. We could pick out trained musicians, but then the study is limited to musicians technically. I guess we would just have to do the vetting procedure and only select the upper echelon of people which can tell extremely minute differences in tempo after a select period of time.

Thoughts I had while working out, lol.

Part II

Papers (Deprecated)

0.0.0.0.1 Preamble. My initial method for retaining information from papers was to just write notes explicitly out of the paper. While this helped to facilitate my understanding, grouping the information in this manner seems to obfuscate any actual knowledge. This method shall be deprecated henceforth, and any information gained from papers will instead be categorized by topic and contributions to that topic.

Chapter 1

Algorithms and Theory

1.1 Monte Carlo and Las Vegas Randomized Algorithms for Systems and Control*: An Introduction [25]

Abstract. *In this paper, we present an introduction to Monte Carlo and Las Vegas randomized algorithms for systems and control. Specific applications of these algorithms include stability analysis, Lyapunov functions, and distributed consensus problems.*

1.2 Approximation Algorithms for k-Anonymity [1]

Abstract. *We consider the problem of releasing a table containing personal records, while ensuring individual privacy and maintaining data integrity to the extent possible. One of the techniques proposed in the literature is k -anonymization. A release is considered k -anonymous if the information corresponding to any individual in the release cannot be distinguished from that of at least $k-1$ other individuals whose information also appears in the release. In order to achieve k -anonymization, some of the entries of the table are either suppressed or generalized (e.g. an Age value of 23 could be changed to the Age range 20-25). The goal is to lose as little information as possible while ensuring that the release is k -anonymous. This optimization problem is referred to as the k -Anonymity problem. We show that the k -Anonymity problem is NP-hard even when the attribute values are ternary and we are allowed only to suppress entries. On the positive side, we provide an $O(k)$ -approximation algorithm for the problem. We also give improved positive results for the interesting cases with specific values of k — in particular, we give a 1.5-approximation algorithm for the special case of 2-Anonymity, and a 2-approximation algorithm for 3-Anonymity.*

1.3 Improved Approximation Algorithms for Maximum Cut and Satisfiability Problems Using Semidefinite Programming [8]

The initial problem attempted to be solved by this paper is creating an approximation algorithm for MAX-CUT to improve on the previous guarantee of $\frac{3}{4}$ for an undirected graph $G = (V, E)$ where $|V| = n$ with nonnegative weights $W_{ij} = W_{ji}$ i.e. \mathbf{W} denotes a weighted adjacency matrix.

We construct a cut of G by assigning to each $i \in V$ a value $y_i \in \{-1, 1\}$ and construct $S \subset G$ by taking $S = \{i | y_i = 1\}$. Define the value of a cut as

$$w(S, \bar{S}) = \frac{1}{2} \sum_{i < j} W_{ij}(1 - y_i y_j) \quad (1.1)$$

The MAX-CUT integer quadratic program is then originally defined as

$$\text{MAX CUT}(G) = \max_{S \subset V} w(S, \bar{S}), \quad (1.2)$$

i.e. the goal is to find values y_i for all $i \in V$ that maximize this equation.

The proposed relaxation is to substitute each y_i in Equation 1.1 with some vector $\mathbf{v}_i \in \mathbb{R}^m$ ($m \leq n$) such that $\|\mathbf{v}_i\|^2 = 1$ so that each \mathbf{v}_i belongs to the m -dimensional unit sphere S_m . Under this relaxation, we instead construct $S \subset G$ by selecting some $\mathbf{r} \in S_m$ randomly and uniformly and take $S = \{i | \langle \mathbf{v}_i, \mathbf{r} \rangle \geq 0\}$ i.e. \mathbf{r} determines a hyperplane which separates vertices of G . Then Equation 1.1 becomes

$$w(S, \bar{S}) = \frac{1}{2} \sum_{i < j} W_{ij}(1 - \langle \mathbf{v}_i, \mathbf{v}_j \rangle). \quad (1.3)$$

Remark. For a real, symmetric matrix A , the following are equivalent:

- A is positive semidefinite;
- $\lambda(A) \subset \mathbb{R}^+ \cup \{0\}$ where $\lambda(X)$ is the set of X 's eigenvalues; and
- $A = B^\top B$ for some $B \in \mathbb{R}^{m \times n}$ where $m \leq n$.

The goal is now to determine optimal values for the unit vectors $\mathbf{v}_1, \dots, \mathbf{v}_n$ and the dimensionality m for the vector space they reside. We can take $\mathbf{B} = [\mathbf{v}_1 \ \dots \ \mathbf{v}_n]$ so that $\mathbf{Y} = \mathbf{B}^\top \mathbf{B}$ is the Gram matrix of \mathbf{B} . Instead of solving for \mathbf{B} right away, we can instead note that each $Y_{ij} = \langle \mathbf{v}_i, \mathbf{v}_j \rangle$ and that our previous constraint is now just the constraint $Y_{ii} = 1$. We also have that \mathbf{Y} is symmetric and PSD by construction. Semidefinite programming can be used to maximize the equation

$$\frac{1}{2} \sum_{i < j} W_{ij}(1 - Y_{ij}). \quad (1.4)$$

Rather than deal with each v_i directly, we can just use an approximate semidefinite programming algorithm to find an optimal \mathbf{Y} in $O\left(\sqrt{n} \left(\log\left(\sum_{i < j} W_{ij}\right) + \log 1/\epsilon\right)\right)$

iterations (Alizadeh's adaptation of Ye's interior-point algorithm) with each iteration taking $O(n^3)$, where $\epsilon > 0$ denotes the acceptable distance from the optimal value Z_P^* (i.e. $Z_P^* - \epsilon$). Once an optimal \mathbf{Y} has been found, an incomplete Cholesky decomposition can be used to obtain $\mathbf{v}_1, \dots, \mathbf{v}_n \in S_m$ s.t.

$$\frac{1}{2} \sum_{i < j} W_{ij} (1 - \langle \mathbf{v}_i, \mathbf{v}_j \rangle) \geq Z_P^* - \epsilon. \quad (1.5)$$

To find m , they note that any *extreme solution*, one which cannot be expressed as the strict convex combination of other feasible solutions, has at most rank l where

$$l \leq \frac{\sqrt{8n+1}-1}{2} < \sqrt{2n}. \quad (1.6)$$

Any other solution has $m \leq n$.

1.3.1 Analysis

The analysis breaks apart into three cases:

1. the general case with nonnegative edge weights,
2. an tighter bound when the weight of the cut constitutes a large portion of the overall weight,
3. and extension to negative weights.

The general case constitutes defining the overall expected weight as in

$$\mathbb{E}[W] = \frac{1}{\pi} \sum_{i < j} W_{ij} \Pr[\operatorname{sgn}(\langle \mathbf{v}_i, \mathbf{r} \rangle) \neq \operatorname{sgn}(\langle \mathbf{v}_j, \mathbf{r} \rangle)]. \quad (1.7)$$

They then state the following lemma:

Lemma 1.1.

$$\Pr[\operatorname{sgn}(\langle \mathbf{v}_i, \mathbf{r} \rangle) \neq \operatorname{sgn}(\langle \mathbf{v}_j, \mathbf{r} \rangle)] = \frac{1}{\pi} \arccos(\langle \mathbf{v}_i, \mathbf{v}_j \rangle), \quad (1.8)$$

which follows from the fact that the probability that these two values have the same sign is equivalent to the dihedral angle between the vectors, i.e. the probability that a randomly selected hyperplane bisects the two divided by all the possible hyperplanes that bisect the unit sphere. This results in the theorem

Theorem 1.1.

$$\mathbb{E}[W] = \frac{1}{\pi} \sum_{i < j} W_{ij} \arccos(\langle \mathbf{v}_i, \mathbf{v}_j \rangle). \quad (1.9)$$

Let $\theta = \arccos(\langle \mathbf{v}_i, \mathbf{v}_j \rangle)$ and define

$$\alpha = \min_{0 \leq \theta \leq \pi} \frac{2}{\pi} \frac{\theta}{1 - \cos \theta}. \quad (1.10)$$

They then provide the following lemmas

Lemma 1.2. For $-1 \leq y \leq 1$, $\arccos(y)/\pi \geq \alpha \cdot \frac{1}{2}(1 - y)$,

which follows from substituting $y = \cos \theta$, and

Lemma 1.3. $\alpha > 0.87856$

which follows from finding a lower bound for Equation 1.10. They use these to show that

Theorem 1.2.

$$\mathbb{E}[W] \geq \alpha \frac{1}{2} \sum_{i < j} W_{ij} (1 - \langle \mathbf{v}_i, \mathbf{v}_j \rangle). \quad (1.11)$$

For the case where the cut constitutes a large portion of the overall weight, we define $W_{\text{tot}} = \sum_{i < j} W_{ij}$, $h(t) = \arccos(1 - 2t)/\pi$, and $\gamma = \arg \min_{0 < t \leq 1} h(t)/t \approx 0.84458$. The analysis then provides the theorem:

Theorem 1.3. Let

$$A = \frac{1}{W_{\text{tot}}} \frac{1}{2} \sum_{i < j} W_{ij} (1 - \langle \mathbf{v}_i, \mathbf{v}_j \rangle). \quad (1.12)$$

If $A \geq \gamma$, then

$$\mathbb{E}[W] \geq \frac{h(A)}{A} \frac{1}{2} \sum_{i < j} W_{ij} (1 - \langle \mathbf{v}_i, \mathbf{v}_j \rangle), \quad (1.13)$$

i.e. as A varies between γ and 1, $h(A)/A$ varies between α and 1. This provides us with a performance guarantee of $h(A)/A - \epsilon$.

1.4 Treewidth, partial k -trees, and chordal graphs [10]

Definition 1.1. A *tree decomposition* of a graph $G = (V, E)$ is a pair

$$(\{X_i \mid i \in I\}, T = (I, M)) \quad (1.14)$$

where $\{X_i \mid i \in I\}$ is a collection of subsets V (also called *bags*), and T is a tree s.t.

- $\cup_{i \in I} X_i = V$ (combining the vertices of all bags forms V)
- $(u, v) \in E \implies \exists i \in I : u, v \in X_i$ (if there is an edge connecting two vertices, then those vertices share at least one bag)
- $\forall v \in V : \{i \in I \mid v \in X_i\}$ induces a connected subtree of T (all bags involving a vertex v form a connected subtree).

Similarly, a *path decomposition* is a tree decomposition such that T is a path (no branches).

Remark. The last statement in Definition 1.1 can be replaced with

- $i, k, j \in I$ and j is on the path from i to k in $T \implies X_i \cap X_k \subseteq X_j$.

Lemma 1.4. [5] Let $(\{X_i \mid i \in I\}, T = (I, M))$ be a tree decomposition of $G = (V, E)$, and let $K \subseteq V$ be a clique in G . Then $\exists i \in I : K \subseteq X_i$.

Definition 1.2. The *width* of a decomposition is $\max_{i \in I} |X_i| - 1$. The *tree-width* ($tw(\cdot)$) is the minimum width over all tree decompositions of G . Likewise, the *path-width* ($pw(\cdot)$) is the minimum width over all path decompositions of G .

Corollary 1.1. For every graph, $tw(G) \geq \omega(G) - 1$ ($\omega(G)$ is the size of the max clique in G).

Theorem 1.4. [2] The following problems are NP-complete:

- Given a graph $G = (V, E)$ and an integer $k < |V|$, is the $tw(G) \leq k$?
- Given a graph $G = (V, E)$ and an integer $k < |V|$, is the $pw(G) \leq k$?

Definition 1.3. The class of K -trees is defined recursively as follows:

- K_{k+1} is a k -tree;
- a k -tree G with $n + 1$ vertices ($n \geq k + 1$) can be constructed from a k -tree H with n vertices by adding a vertex adjacent to exactly k vertices, namely all vertices of a $(k + 1)$ -clique of H .

Theorem 1.5. [22] Let $G = (V, E)$ be a graph. The following statements are equivalent:

- G is a k -tree.
- G is a connected, G has a k -clique, but no $(k + 2)$ -cliques, and every minimal separator of G is a k -clique.
- G is connected, $|E| = k|V| - \frac{1}{2}k(k + 1)$, and every minimal separator of G is a k -clique.
- G has a k -clique, but not a $(k + 2)$ -clique, and every minimal separator of G is a clique, and for all distinct non-adjacent pairs of vertices $x, y \in V$, there exist exactly k vertex disjoint paths from x to y .

Definition 1.4. A partial k -tree is a grpah that contains all the vertices and a subset of the edges of a k -tree.

Theorem 1.6. [16] G is a partial k -tree if and only if G has tree-width at most k .

Chapter 2

Deep Learning

2.1 Efficient Per-Example Gradient Computations [9]

An important note about this paper is that it is *not* a derivation of the per-example gradient itself. Rather, it “describes an efficient technique for computing the *norm* of the gradient” with respect to the loss function.

Let each layer of a neural network be defined by the standard transformations

$$\begin{aligned}\mathbf{z}^{(i)} &= \mathbf{W}^{(i)} \mathbf{h}^{(i-1)} \\ \mathbf{h}^{(i)} &= \phi^{(i)}(\mathbf{z}^{(i)})\end{aligned}\tag{2.1}$$

with the bias assumed to be an extra row/column of \mathbf{W} and loss function \mathcal{L} . Let $\mathcal{B} \subset \mathcal{D}$ be a minibatch of the input dataset \mathcal{D} , and let $\mathcal{B}^{[j]}$ be the j th example in the minibatch with $\mathcal{L}^{[j]}$ corresponding to the loss of $\mathcal{B}^{[j]}$ and cost function $C = \sum_{j=1}^n \mathcal{L}^{[j]}$. The per example gradient norm is computed with respect both to the example and the layer, such that the gradient norm for layer i and example j is

$$\left\| \frac{d\mathcal{L}^{[j]}}{d\mathbf{W}^{(i)}} \right\|^2 = \sum_{k,l} \left(\frac{\partial \mathcal{L}^{[j]}}{\partial W_{k,l}^{(i)}} \right)^2,\tag{2.2}$$

which is just the frobenius norm. Note that this equation can be used to calculate the L_2 norm of the parameter gradient for an example or the L_2 norm of the gradient for an individual weight matrix by summing over j or i , respectively.

The proposed method makes use of $\mathbf{H}^{(i)}$ whose j th row contains the activation layer $\mathbf{h}^{(i)}$ corresponding to $\mathcal{B}^{[j]}$ i.e. $\mathbf{H}_j^{(i)} = \phi^{(i)}(\mathbf{W}^{(i)} \phi^{(i-1)}(\dots(\mathbf{W}^{(1)} \mathcal{B}^{[j]})\dots))$. Likewise, define $\mathbf{Z}^{(i)}$ for each $\mathbf{z}^{(i)}$ and compute $\bar{\mathbf{Z}} = \nabla_{\mathbf{Z}} C$, which can be computed in a single pass using standard backpropagation. This gives us

$$\left\| \frac{\partial \mathcal{L}^{[j]}}{\partial \mathbf{W}^{(i)}} \right\|^2 = \left(\sum_k (\bar{Z}_{j,k}^{(i)})^2 \right) \left(\sum_k (H_{j,k}^{(i-1)})^2 \right) = \left\| \bar{\mathbf{Z}}_j^{(i)} \right\|^2 \left\| \mathbf{H}_j^{(i-1)} \right\|^2.\tag{2.3}$$

2.1.1 Commentary

This method works well with gradient clipping due to the efficient computation of gradient norms. This calculation allows for simple re-scaling of the gradients used in backwards propagation.

2.2 Understanding gradient clipping in incremental gradient methods [21]

This paper considers the class of problems dealing with the minimization of

$$\min_{x \in \mathbb{R}^n} f(x) \quad \text{where} \quad f(x) = \frac{1}{m} \sum_{i=1}^m f_i(x), \quad (2.4)$$

where $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ are differentiable, non-convex functions. As a side note, this differs from notation that I'm used to, in that x is considered to be the set of parameters while i , or more importantly f_i , is a function of the parameters using the i^{th} component of the dataset to measure the loss. That is, the function $f_i(x)$ determines the loss accumulated when using x as the set of parameters on input i . f is said to be the *objective function* while each f_i is called a *component function*.

The incremental method for SGD is

$$x_{k+1} = x_k - \alpha_k \nabla f_{i_k}(x_k) \quad k = 0, 1, \dots \quad (2.5)$$

where x_k is the values of the parameters at iteration k , α_k is the learning rate, and f_{i_k} is a uniform randomly chosen component function $i_k \in \{1, 2, \dots, m\}$ of the dataset. The incremental method for IGC is

$$x_{k,i} = x_{k,i-1} - \alpha_k \nabla f_i(x_{k,i-1}) \quad i = 1, \dots, m \quad k = 0, 1, \dots \quad (2.6)$$

where the recursion is defined by base case $x_{0,0}$ and iterative step $x_{k,0} = x_{k-1,m}$ and is more cyclical by nature.

The clipping function $\mathcal{C} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ with clipping threshold parameter $\eta > 0$ is defined by

$$\mathcal{C}(g; \eta) = \min \left\{ 1, \frac{\eta}{\|g\|} \right\} \cdot g. \quad (2.7)$$

Our iterative methods can take advantage of clipping by clipping the gradient so that SGD with clipping is

$$x_{k+1} = x_k - \alpha_k \mathcal{C}(\nabla f_{i_k}(x_k); \eta), \quad (2.8)$$

and IGC with clipping is

$$x_{k,i} = x_{k,i-1} - \alpha_k \mathcal{C}(\nabla f_i(x_{k,i-1}); \eta), \quad (2.9)$$

and is written iteratively (using the notation $x_k = x_{k,0} = x_{k-1,m}$) as

$$x_{k+1} = x_k - \alpha_k \sum_{i=1}^m \mathcal{C}(\nabla f_i(x_{k,i-1}); \eta), \quad (2.10)$$

For the theoretical understanding of the clipping function's effects on the convergence of both SGD and IGC, we make the following assumptions:

Bounded below there exists an f^* such that for all $x \in \mathbb{R}^n$, $f(x) \geq f^*$;

Relaxed smoothness there exist constants $L_0, L_1 \in \mathbb{R}^+$ such that $\|\nabla^2 f(x)\| \leq L_0 + L_1 \|\nabla f(x)\|$.

The second condition becomes Lipschitz smoothness when $L_1 = 0$. Let $f_k = f(x_k)$ and assume for simplicity from here on out that SGD and IGC have a constant step size α and that IGC's notation for $x_{k,i}$ can be simplified for $x_{k,0}$ to x_k . The paper shows that convergence for the algorithm can be ensured if there exists a constant C such that

$$\left\langle \sum_{i=1}^m \mathcal{C}(\nabla f_i(x_k); \eta), \nabla f_k \right\rangle \geq C \|\nabla f_k\|^2 \quad (2.11)$$

i.e. if the magnitude of the clipped gradient projected in the direction of the regular gradient is bounded below by the magnitude of the regular gradient. This value is simplified further as follows.

Let $g_{ki} = \nabla f_i(x_k)$ and note the following properties of g_{ki}

$$\begin{aligned} \beta_{ki} &= \frac{\langle g_{ki}, \nabla f_k \rangle}{\|\nabla f_k\|^2} = \frac{\|g_{ki}\|}{\|\nabla f_k\|} \cos \theta_{ki} \\ g_{ki} &= \beta_{ki} \nabla f_k + t_{ki} \quad \text{where} \end{aligned} \quad . \quad (2.12)$$

$$t_{ki} \in \nabla f_k^\perp = \{z | \langle z, \nabla f_k \rangle = 0\}$$

Since $\frac{1}{m} \sum_i^m g_{ki} = \nabla f_k$, it is a necessary condition that $\sum_i^m \beta_{ki} = m$ and $\sum_i^m t_{ki} = 0$. Let $\gamma_{ki} = \min\{1, \eta/\|g_{ki}\|\}$ and note that $\mathcal{C}(g_{ki}; \eta) = \gamma_{ki} g_{ki}$. The paper then shows that

$$\left\langle \sum_{i=1}^m \mathcal{C}(\nabla f_i(x_k); \eta), \nabla f_k \right\rangle = \sum_{i=1}^m \gamma_{ki} \beta_{ki} \|\nabla f_k\|^2. \quad (2.13)$$

so that Equation 2.11 becomes

$$\sum_{i=1}^m \gamma_{ki} \beta_{ki} \geq C, \quad (2.14)$$

that is, *the necessary condition for convergence is that the sum in Equation 2.14 is greater than some positive scalar C .*

There are then essentially two conditions, either of which will ensure the convergence of our algorithm:

1. all g_{ki} can be partitioned into pairs so that if g_{kj} and g_{kl} are a pair then $\beta_{kj} + \beta_{kl} > 0$ and $\|t_{kj}\| = \|t_{kl}\|$;
2. all $\beta_{ki} > 0$ or equivalently all $\cos \theta_{ki} > 0$.

Under the second condition, the paper then shows that a lower bound can be derived making use of the minimum and maximum gradient magnitudes w.r.t. an example i . That is, if we let $m_g = \min_i \|g_{ki}\|$ and $M_g = \max_i \|g_{ki}\|$, and we let constants c_1 and c_2 be defined by the inequalities

$$0 < c_1 \leq \frac{1}{m} \sum_{i=1}^m \cos \theta_{ki} \quad \text{and} \quad 0 < c_2 \leq \frac{m_g}{M_g}, \quad (2.15)$$

then the lower bounds for Equation 2.11 are then divided into three scenarios:

1. if $\eta \geq M_g$, then

$$\left\langle \sum_{i=1}^m \mathcal{C}(\nabla f_i(x_k); \eta), \nabla f_k \right\rangle = m \|\nabla f_k\|^2; \quad (2.16)$$

2. if $\eta \leq m_g$, then

$$\left\langle \sum_{i=1}^m \mathcal{C}(\nabla f_i(x_k); \eta), \nabla f_k \right\rangle = \left(\eta \sum_{i=1}^m \cos \theta_{ki} \right) \|\nabla f_k\|; \quad (2.17)$$

3. and if $m_g < \eta < M_g$, then

$$\left\langle \sum_{i=1}^m \mathcal{C}(\nabla f_i(x_k); \eta), \nabla f_k \right\rangle \geq m \min\{\eta c_1, c_2 \|\nabla f_k\|\} \|\nabla f_k\|. \quad (2.18)$$

The paper concludes by providing the following theorems on the behavior of SGD and IGC with gradient clipping.

Theorem 2.1. *Assuming that we have a lower bound and relaxed smoothness, if*

$$\alpha \leq \frac{c_1}{4\eta L_1} \quad (2.19)$$

then the iterates $\{x_k\}$ generated by Equation 2.8 satisfy

$$\frac{1}{T+1} \sum_{k=0}^T \min\{\eta c_1, c_2 \|\nabla f_k\|\} \|\nabla f_k\| \leq \frac{2(f_0 - f^*)}{(T+1)\alpha} + 5\alpha\eta^2 L_0 + \frac{\alpha\eta^3 L_1 c_1}{c_2}. \quad (2.20)$$

Theorem 2.2. *Assuming that we have a lower bound, relaxed smoothness, and each component function also satisfies relaxed smoothness i.e.*

$$\|\nabla^2 f_i(x)\| \leq L_{0i} + L_{1i} \|\nabla f_i(x)\|, \quad i = 1, \dots, m; \quad (2.21)$$

if

$$\alpha \leq \min \left\{ \frac{c_1}{2m(2\eta L_1 + G)}, \frac{1}{5mL_{0i}}, \frac{1}{8m\eta L_1 i} \right\}, \quad i = 1, \dots, m; \quad (2.22)$$

then the iterates $\{x_k\}$ generated by Equation 2.10 satisfy

$$\frac{1}{T+1} \sum_{k=0}^T \min\{\eta c_1, c_2 \|\nabla f_k\|\} \|\nabla f_k\| \leq \frac{2(f_0 - f^*)}{(T+1)\alpha m} + 5\alpha m \eta^2 L_0 + \frac{\alpha m \eta^2 c_1 (2\eta L_1 + G)}{2c_2}, \quad (2.23)$$

where

$$G = 5 \max_{1 \leq i \leq m} L_{0i} + 16\eta \max_{1 \leq i \leq m} L_{1i}. \quad (2.24)$$

2.3 Pruning neural networks without any data by iteratively conserving synaptic flow [24]

This paper, while theoretically captivating, is lackluster with providing mathematical commentary on the concepts and proofs. It may be worthwhile to go back and either note the more mathematical definitions of the paper, or derive my own definitions and reprove their claims using a more solid theoretical foundation. This may even set a better groundwork for future analysis.

2.3.1 Pruning Algorithms

Pruning algorithms are generally defined by

1. scoring parameters by some metric, and
2. masking parameters according to their scores.

The latter is generally done by removing the parameters (e.g. making the value zero via Hadamard product). This can either be done via global masking or layer-masking, with global-masking performing better but suffering from *layer-collapse*, which is when an entire layer is masked.

Let θ be a collection of network parameters and θ_{prune} be the parameters remaining after pruning. The *compression ratio* ρ of a pruning algorithm is

$$\rho = \frac{|\theta|}{|\theta_{\text{prune}}|} \quad (2.25)$$

. We define ρ_{\max} as the maximal possible compression ratio for a *network* that doesn't lead to layer collapse and the ρ_{cr} as the maximal compression ratio a given *algorithm* can achieve without inducing layer collapse. Note that the distinction in the two is ρ_{\max} is maximal for a network while ρ_{cr} is maximal for an algorithm. These definitions motivate the following axiom:

Axiom 2.1 (Maximal Critical Compression). *For any pruning algorithm and any network, we should always have $\rho_{\text{cr}} = \rho_{\max}$.*

2.3.2 Synaptic Saliency

Synaptic saliency is a class of score metrics defined by

$$\mathcal{S}(\theta) = \frac{\partial \mathcal{R}}{\partial \theta} \odot \theta \quad (2.26)$$

where \mathcal{R} is a scalar loss function of the output y of a feed-forward network parameterized by θ . This metric satisfies two conservation laws:

Theorem 2.3 (Neuron-wise Conservation of Synaptic Saliency). *For a feedforward neural network with continuous, homogeneous activation functions $\phi(x) = \phi'(x)x$, let j be the index of a hidden neuron in layer i . The sum of the synaptic saliency for the incoming parameters to a hidden neuron $\left(\mathcal{S}_j^{(i)_{in}} = \left\langle \frac{\partial \mathcal{R}}{\partial \theta_j^{(i)}}, \theta_j^{(i)} \right\rangle \right)$ is equal*

to the sum of the synaptic saliency for the outgoing parameters from the hidden neuron $\left(\mathcal{S}_j^{(l)out} = \langle \frac{\partial \mathcal{R}}{\partial \theta_{:,j}^{(l+1)}}, \theta_{:,j}^{(l+1)} \rangle \right)$.

Theorem 2.4 (Network-wise Conservation of Synaptic Saliency). *The sum of the synaptic saliency across any set of parameters that exactly separates the input neurons x from the output neurons y of a feedforward neural network with homogeneous activation functions equals $\langle \frac{\partial \mathcal{R}}{\partial x}, x \rangle = \langle \frac{\partial \mathcal{R}}{\partial y}, y \rangle$.*

2.3.3 Algorithm

For the following theorem, we define the *prune size* to be the total score for the parameters pruned at any iteration and the *cut size* to be the total score for an entire layer.

Theorem 2.5. *If a pruning algorithm with global-masking assigns positive scores that respect layer-wise conservation, and if the prune size is strictly less than the cut size whenever possible, then the algorithm satisfies the Maximal Critical Compression axiom.*

The pruning algorithm defined in this paper uses a loss function

$$\mathcal{R}_{SF} = \mathbf{1}^\top \left(\prod_{l=1}^L |\theta^{(l)}| \right) \mathbf{1} \quad (2.27)$$

so that for a fully connected network ($f_\theta(x) = \theta^{(L)} \dots \theta^{(1)} x$) the Synaptic Flow score for a parameter $\theta_{i,j}^{(l)}$ is

$$\mathcal{S}_{SF}(\theta_{i,j}^{(l)}) = \left[\mathbf{1}^\top \prod_{k=l+1}^L |\theta^{(k)}| \right]_i |\theta_{i,j}^{(l)}| \left[\prod_{k=1}^{l-1} |\theta^{(k)}| \mathbf{1} \right]_j \quad (2.28)$$

i.e. the portion of the l_1 -path norm the network has through this parameter. This contributes to the development of Algorithm 1.

Algorithm 1 Iterative Synaptic Flow Pruning (SynFlow)

Require: network f_θ , compression ratio ρ , iteration steps n

```

 $\mu = \mathbf{1}$  ▷ Initialize binary mask
for  $k$  in  $[1, \dots, n]$  do
     $\theta_\mu \leftarrow \mu \odot \theta_0$  ▷ Mask parameters
     $\mathcal{R} \leftarrow \mathbf{1}^\top \left( \prod_{l=1}^L |\theta_\mu^{(l)}| \right) \mathbf{1}$  ▷ Evaluate SynFlow objective
     $\mathcal{S} = \frac{\partial \mathcal{R}}{\partial \theta_\mu} \odot \theta_\mu$  ▷ Compute SynFlow score
     $\tau \leftarrow (1 - \rho^{-k/n})$  percentile of  $\mathcal{S}$  ▷ Find threshold
     $\mu \leftarrow (\tau < \mathcal{S})$  ▷ Update mask
end for
return  $f_{\mu \odot \theta}$  ▷ Return masked network

```

2.3.4 Commentary

A thought I had: the loss function defined in Equation 2.27 seems to almost indicate a just-past-the-minimum metric for the predefined network. It effectively is just the sum of the entries in the product of the network space. Although I agree with the data agnostic approach, I feel there should be a more suitable loss function which should seek to preserve this score rather than minimize it, or at the very least, should compare this score against some optimum, i.e. the optimum score for that network under those parameters. I also feel that the prune size and cut size should have a greater impact over the outcome of the resulting network, such that the compression ratio is fairly equal for all layers.

Research Question 2.1. Does SynFlow, or other pruning methods for that matter, produce the same network given different initial parameterizations? Furthermore, might we be able to use a genetic algorithm to find the “lottery ticket”?

One thing that needs to be considered with this question is that it’s entirely possible that isomorphic networks could be produced, so there may need to be some kind of check that occurs post pruning that confirms whether or not this network has been discovered already.

2.4 Learning Low-rank Deep Neural Networks via Singular Vector Orthogonality Regularization and Singular Value Sparsification [28]

This paper focuses on an improvement on Deep Neural Networks (specifically ones involving convolutional layers) by utilizing the SVD of a convolution and sparsifying the underlying singular values. This is done by representing the kernel \mathbf{K} of a convolution as a 4-D tensor with dimension $\mathbf{K} \in \mathbb{R}^{n \times c \times w \times h}$ where n is the number of filters (output channels), c is the number of input channels, and w, h are the width and height of the kernel respectively. They break the kernel down into the *channel-wise decomposition* and the *spatial-wise decomposition*.

Under channel-wise decomposition, \mathbf{K} is reshaped into a 2-D matrix $\hat{\mathbf{K}} \in \mathbb{R}^{n \times cwh}$, which can be decomposed with SVD into $\mathbf{U} \in \mathbb{R}^{n \times r}$, $\mathbf{V} \in \mathbb{R}^{cwh \times r}$, and $\mathbf{s} \in \mathbb{R}^r$ where \mathbf{U}, \mathbf{V} are unitary and $r = \min(n, cwh)$. The convolution then becomes the multiplication of convolutions $\mathbf{K}_1 \in \mathbb{R}^{r \times c \times w \times h}$ (reshaped from $\text{diag}(\sqrt{\mathbf{s}}) \mathbf{V}^\top$) and $\mathbf{K}_2 \in \mathbb{R}^{c \times r \times 1 \times 1}$ (reshaped from $\mathbf{U} \text{diag}(\sqrt{\mathbf{s}})$), a simplification discovered by [31].

Under spatial-wise decomposition, \mathbf{K} is reshaped into a 2-D matrix $\hat{\mathbf{K}} \in \mathbb{R}^{nw \times ch}$, which can be decomposed with SVD into $\mathbf{U} \in \mathbb{R}^{nw \times r}$, $\mathbf{V} \in \mathbb{R}^{ch \times r}$, and $\mathbf{s} \in \mathbb{R}^r$ where \mathbf{U}, \mathbf{V} are unitary and $r = \min(n, cwh)$. The convolution then becomes the multiplication of convolutions $\mathbf{K}_1 \in \mathbb{R}^{r \times c \times 1 \times h}$ (reshaped from $\text{diag}(\sqrt{\mathbf{s}}) \mathbf{V}^\top$) and $\mathbf{K}_2 \in \mathbb{R}^{c \times r \times 1 \times 1}$ (reshaped from $\mathbf{U} \text{diag}(\sqrt{\mathbf{s}})$).

In either case, the SVD decomposition of the matrix is stored in memory rather than the convolution, and the convolution is reconstructed from these matrices.

Since the assurance that \mathbf{U}, \mathbf{V} are orthogonal is a costly procedure, an orthogonality regularizer is added to the cost of the loss function. This orthogonality

Figure 2.1: Effect of applying different sparsity-inducing regularizers during SVD training. All models are achieved with Spatial-wise decomposition.

regularizer is given by

$$L_o(\mathbf{U}, \mathbf{V}) = \frac{1}{r^2} \left(\|\mathbf{U}^\top \mathbf{U} - \mathbf{I}\|_F^2 + \|\mathbf{V}^\top \mathbf{V} - \mathbf{I}\|_F^2 \right). \quad (2.29)$$

This assures that the matrices will be approximately semi-unitary. Note that they can not be orthogonal since they are not square.

As part of the sparsification, another regularizer is used for sparsifying the singular values. This regularizer was tested using both the L^1 norm and a *Hoyer regularizer*, given by

$$L^H(\mathbf{s}) = \frac{\|\mathbf{s}\|_1}{\|\mathbf{s}\|_2}. \quad (2.30)$$

The results given by Figure 2.4 show that while the L^1 regularizer doesn't match the performance/accuracy tradeoff of the hoyer regularizer, allowing for a higher accuracy loss can result in better performance by the L^1 regularizer. I.e. extremely large compression rates which result in higher accuracy loss can provide a more generous reduction in #FLOPs under L^1 , whereas having a more moderate compression and a lower accuracy loss has better reduction of #FLOPs under L^H .

The overall objective function is finally given by

$$L(\mathbf{U}, \mathbf{s}, \mathbf{V}) = L_T(\text{diag}(\sqrt{\mathbf{s}}) \mathbf{V}^\top, \mathbf{U} \text{diag}(\sqrt{\mathbf{s}})) + \lambda_o \sum_{l=1}^D L_o(\mathbf{U}^{(l)}, \mathbf{V}^{(l)}) + \lambda_s \sum_{l=1}^D L_s(\mathbf{s}^{(l)}), \quad (2.31)$$

where L_T is the training loss computed with decomposed layers, L_o is the orthogonality loss from Equation 8.7, L_s is the sparsity-inducing regularization loss (either L^1 or L^H), and λ_o, λ_s are hyperparameters for controlling the effect of the objective functions on the overall loss. The paper demonstrates that the change in accuracy from changing $\lambda_o = 1.0$ to $\lambda_o = 0.0$ is about -2.0% , showing that the orthogonality causes a definite change in performance.

As for pruning, an energy threshold $e \in [0, 1]$ is a hyperparameter which for each layer, a set \mathbb{K} is determined containing the largest number of singular values s.t.

$$\sum_{j \in \mathbb{K}} s_j^2 \leq e \sum_{i=1}^r s_i^2. \quad (2.32)$$

The values in \mathbb{K} are then pruned by setting each $s_j = 0$.

TODO: Fill in some details on channel-wise and spatial-wise decomposition because I have no idea what either of these things are.

2.5 Scalable Adaptive Stochastic Optimization Using Random Projections [15]

This paper uses a similar notation to that of Chapter 2.2 in that we consider the function

$$F_t = f_t(\beta) + \phi(\beta) \quad (2.33)$$

where f_t is a convex loss function, ϕ is a regularizer, and $\beta \in \mathbb{R}^p$ is the set of parameters.

Definition 2.1. The *regret* after T rounds is defined as

$$R(T) = \sum_{t=1}^T F_t(\beta_t) - \sum_{t=1}^T F_t(\beta^{\text{opt}}), \quad (2.34)$$

where β_t is the set of parameters at round t and β^{opt} is a fixed optimal predictor.

The paper proceeds to define $\mathbf{g}_t \in \nabla f_t(\beta_t)$ as a *subgradient* of the loss function, which we may view as a gradient obtained from a minibatch of our dataset. Standard, first-order methods use a fixed learning rate η to move in the opposite direction of \mathbf{g}_t , but their method builds on second-order methods which use an adaptive learning rate which is controlled by a time-varying *proximal term*, i.e. a function determining the adaptive learning rate. Given $\mathbf{G}_t = \sum_{i=1}^t \mathbf{g}_i \mathbf{g}_i^\top$ and $\mathbf{H}_t = \delta \mathbf{I}_p + \mathbf{G}_t^{1/2}$ The proximal terms for ADAFULL and ADAGRAD are given by

$$\psi_t(\beta) = \frac{1}{2} \langle \beta, \mathbf{H}_t \beta \rangle \quad \text{and} \quad \psi_t(\beta) = \frac{1}{2} \langle \beta, (\delta \mathbf{I}_p + \text{diag}(\mathbf{G}_t)^{1/2}) \beta \rangle \quad (2.35)$$

respectively. When p is large, ADAFULL may be computationally infeasible, yet ADAGRAD is unable to capture dependencies between coordinates in gradient terms.

Definition 2.2. *Effective rank* is defined for a matrix Σ as

$$r(\Sigma) = \frac{\text{tr}(\Sigma)}{\|\Sigma\|}. \quad (2.36)$$

Σ is considered to have *low effective rank* when $r(\Sigma) \leq \text{rank}(\Sigma) \leq p$.

Remark. Low effective rank is a common trademark of high-dimensional data, as there is generally a lot of sparsity.

Definition 2.3. *Random projections* are low-dimensional embeddings $\Pi : \mathbb{R}^p \rightarrow \mathbb{R}^\tau$ which preserve - up to a small distortion - the geometry of a subspace of vectors.

Definition 2.4. The *Subsampled Randomized Fourier Transform* is a *structured* random projection given by

$$\Pi = \sqrt{p/\tau} \mathbf{S} \Theta \mathbf{D} \quad (2.37)$$

where

1. $\mathbf{S} \in \mathbb{R}^{\tau \times p}$ is a subsampling matrix,
2. $\mathbf{D} \in \mathbb{R}^{p \times p}$ is a diagonal matrix whose entries are drawn independently from $\{-1, 1\}$,
3. $\Theta \in \mathbb{R}^{p \times p}$ is a unitary discrete fourier transform matrix.

They conclude by proposing Algorithms 2 and 3.

Algorithm 2 ADA-LR

Require: $\eta > 0, \delta \geq 0, \tau$

- 1: **for** $t = 1 \dots T$ **do**
 - 2: Receive $\mathbf{g}_t = \nabla f_t(\beta_t)$
 - 3: $\mathbf{G}_t = \mathbf{G}_{t-1} + \mathbf{g}_t \mathbf{g}_t^\top$ ▷ Construct additive gram matrix
 - 4: Project $\tilde{\mathbf{G}}_t = \mathbf{G}_t \boldsymbol{\Pi}$
 - 5: $\mathbf{Q}\mathbf{R} = \tilde{\mathbf{G}}_t$ ▷ QR Decomposition to get the basis vectors \mathbf{Q}
 - 6: $\mathbf{B} = \mathbf{Q}^\top \tilde{\mathbf{G}}_t$ ▷ Project onto new basis
 - 7: $\mathbf{U}, \Sigma, \mathbf{V} = \mathbf{B}$ ▷ SVD
 - 8: $\beta_{t+1} = \beta_t - \eta \mathbf{V} (\Sigma^{1/2} + \delta \mathbf{I})^{-1} \mathbf{V}^\top \mathbf{g}_t$
 - 9: **end for**
-

Algorithm 3 RADA GRAD

Require: $\eta > 0, \delta \geq 0, \tau$

- 1: **for** $t = 1 \dots T$ **do**
- 2: Receive $\mathbf{g}_t = \nabla f_t(\beta_t)$
- 3: Project $\tilde{\mathbf{g}}_t = \Pi \mathbf{g}_t$
- 4: $\tilde{\mathbf{G}}_t = \tilde{\mathbf{G}}_{t-1} + \mathbf{g}_t \tilde{\mathbf{g}}_t^\top$
- 5: $\mathbf{Q}_t, \mathbf{R}_t \leftarrow \text{qr_update}(\mathbf{Q}_{t-1}, \mathbf{R}_{t-1}, \mathbf{g}_t, \tilde{\mathbf{g}}_t)$
- 6: $\mathbf{B} = \tilde{\mathbf{G}}_t^\top \mathbf{Q}_t$
- 7: $\mathbf{U}, \Sigma, \mathbf{W} = \mathbf{B}$ ▷ SVD
- 8: $\mathbf{V} = \mathbf{W} \mathbf{Q}^\top$
- 9: $\gamma_t = \eta(\mathbf{g}_t - \mathbf{V} \mathbf{V}^\top \mathbf{g}_t)$
- 10: $\beta_{t+1} = \beta_t - \eta \mathbf{V} (\Sigma^{1/2} + \delta \mathbf{I})^{-1} \mathbf{V}^\top \mathbf{g}_t - \gamma_t$
- 11: **end for**

2.6 A Mini-Block Natural Gradient Method for Deep Neural Networks [3]

Definition 2.5. The *empirical fisher matrix* F_θ of a model's conditional distribution is defined as

$$F_\theta = \mathbb{E}_{x \sim X, y \sim p_\theta(\cdot|x)} \left[\frac{\partial \log p_\theta(y|x)}{\partial \theta} \left(\frac{\partial \log p_\theta(y|x)}{\partial \theta} \right)^\top \right] \quad (2.38)$$

where X is a data distribution and $p_\theta(\cdot|x)$ is the density function of the conditional distribution defined by the model with a given input x .

Remark. [19] has shown that when the conditional distribution is a categorical distribution for classification or Gaussian distribution for regression, then the empirical fisher matrix can be written compactly as

$$F_\theta = \frac{1}{n} J_\theta^\top J_\theta \quad (2.39)$$

where $J_\theta \in \mathbb{R}^{n \times p}$ is the Jacobian of the loss function on inputs $x \in X$ with respect to the parameters, i.e. if $\mathcal{L}_\theta^{[i]}$ identifies the loss when evaluating f_θ on example $x^{[i]}$ of the input and $\mathcal{L}_\theta = \mathcal{L}_\theta^{[:]}$, then the Jacobian is given by

$$J_\theta = \frac{\partial \mathcal{L}_\theta}{\partial \theta} \quad (2.40)$$

Since we often express our network a series of linear transformations via matrix multiplication, the Jacobian and empirical Fisher matrices of matrices should be assumed to be vectorized versions of their matrix counterparts, e.g.

$$J_{\mathbf{W}^{(l)}} = \frac{\partial \mathcal{L}_\theta}{\partial \text{vec}(\mathbf{W}^{(l)})}. \quad (2.41)$$

The paper proposes to approximate the empirical fisher matrix, as the actual one is computationally infeasible. This is done via

$$F_{\boldsymbol{\theta}} \approx \text{diag}(F_{\mathbf{W}^{(1)}}, \dots, F_{\mathbf{W}^{(L)}}). \quad (2.42)$$

The pre-conditioning on the gradient is then done as

$$\mathbf{W}_{t+1}^{(l)} \leftarrow \mathbf{W}_t^{(l)} - \alpha F_{\mathbf{W}_t^{(l)}}^{-1} \frac{\partial \mathcal{L}_{\boldsymbol{\theta}}}{\partial \mathbf{W}_t^{(l)}} \quad (2.43)$$

2.6.1 Preconditioning

2.6.1.1 Fully-Connected, Feed Forward Layers

For fully-connected, feed forward networks, this is accomplished for each neuron by constructing the empirical fisher approximation over a subset of the neurons (i.e. subset of the rows) *independently and in parallel* so that if \mathbf{B} is any subset of the rows of $\mathbf{W}^{(l)}$, then the gradient update for this *mini-block* \mathbf{B} is given by

$$\mathbf{B}_{t+1}^{(l)} \leftarrow \mathbf{B}_t^{(l)} - \alpha F_{\mathbf{B}_t^{(l)}}^{-1} \frac{\partial \mathcal{L}_{\boldsymbol{\theta}}}{\partial \mathbf{B}_t^{(l)}} \quad (2.44)$$

and if B_1, \dots, B_n form any partition of the rows of $\mathbf{W}^{(l)}$, then the complete gradient update is performed by computing Equation 2.44 for each B_i in parallel and recombining the results in their original order.

2.6.1.2 Convolutional Layers

TODO: Fill this out

Chapter 3

Finance

3.1 An Introduction to Quantitative Finance [4]

It is important to note that I only present the conclusions to propositions presented in this book rather than deriving the proofs themselves. This is especially the case whenever I list out the value of any derivative: these values are derived in the book using various proofs and are *not* defined to have these values in the first place. For the most part, any derivative with maturity T has a definitive value at maturity, but is a random variable up until that point. The book spends a good amount of time providing deterministic values for these r.v.s, and all that is generally listed here are these deterministic values.

3.1.1 Interest Rates

Definition 3.1. A *notional* or *principal* is an initial deposit or value N .

Definition 3.2. An *interest rate* r is the rate at which a value is increased according to a specified frequency of time.

Suppose we have an account with interest rate r and notional N . The value of the account at time T is

$$N_T = Ne^{rT} \quad (3.1)$$

Throughout finance, there exist discrete and continuous analogs for most equations due the discrete case being more realistic while the continuous case is far more mathematically efficient. That being said, discrete analogs will be used sparsely and we will primarily depend on continuous equations. The discrete versions can be derived from the continuous versions by noting that

$$e^{rT} = \left(1 + \frac{r_m}{m}\right)^{mT} \quad (3.2)$$

where T is commonly denoted as the time in years, m is the frequency the interest is compounded, and r_m is the equivalent interest rate with the discrete frequency. r_m can then be recovered from r using

$$r_m = m \left(e^{\left(\frac{r}{m}\right)} - 1 \right). \quad (3.3)$$

Table 3.1: Daycount conventions and accrual factors

Daycount convention	α for 16 December 2011 - 16 March 2012
act/365	91/365
act/act	15/365 + 76/366
act/360	91/360
30/365	1/4

Definition 3.3. A *money market account* is an account compounded continuously at rate r with notional 1. I.e. $M_0 = 1$ and $M_t = e^{rt}$.

Definition 3.4. A *zero coupon bond* (ZCB) with maturity T is an asset that pays 1 at time T (and nothing else). The value of a ZCB at time t for a continuously compounded rate r is denoted as

$$Z(t, T) = e^{-r(T-t)} \quad (3.4)$$

where $Z(T, T) = 1$ by definition. The values $Z(t, T)$ with $0 \leq t \leq T$ are known as *discount factors* or *present values*.

The intuition behind the present values for a ZCB derives from the fact that if we have two portfolios, one a ZCB with maturity T and another that contains cash which will accumulate interest to be worth 1 at time T , then both will be worth the same value at each t leading up to T and thereafter.

Definition 3.5. An *annuity* a series of fixed cashflows at specified times $T_i = 1, \dots, n$. The value at time t is given by

$$V = C \sum_{i=1}^n Z(t, T_i). \quad (3.5)$$

When we substitute in Equation 3.4, we notice that this is just a geometric series. Because of this, if an annuity is payed once a year for M years with rate r , this equates to

$$V = \sum_{i=1}^n \frac{1}{(1+r)^i} = \frac{1}{r} \left(1 - \frac{1}{(1+r)^M} \right) \quad (3.6)$$

Definition 3.6. The *accrual factor* is (under the discrete analogy) the frequency which interest is compounded, and is denoted as α .

In Equations 3.2 and 3.3, this is the same as taking $\alpha = 1/m$. Due to the year being uneasily divided by simple values of α , several *daycount conventions* have been made for different settings. We list a sample of these in Table 3.1. The daycount conventions are seemingly deceptive in that they hide various details behind terminology. For example, two cases are 30/x and x/360, which respectively denote that each month in the period should be counted as only having 30 days and each month in a year should be counted as only having 30 days. I want to emphasize the weird convention because when we encounter a fraction in the work

such as 30/act, this means to say that if we specify a period of three months, then this fraction actually denotes

$$\frac{30 * 3}{m_1 + m_2 + m_3} \quad (3.7)$$

where each m_i is the number of actual days in the specified month. Like wise, if we have a daycount act/act, this denotes the number specifies the number of actual days in the period divided by the actual days in the year. If we have a leap year, then the denominator contains 366, whereas if February is included in the daycount, then we may have 28 or 29 days dependent on the year. I hope this helps to explain it better, if not then *oof*.

A few specific cases arise but are only defined in the exercises. I will list those here.

Definition 3.7. *Semi-bond* is the US dollar interest rate which is semi-annual compounding with 30/360 daycount and is denoted as γ_{SB} .

Definition 3.8. *Annual money* is the US dollar interest rate which is annual compounding with act/360 daycount and is denoted as γ_{AM} .

If we know the rate r_α and accrual factor α of a quoted interest rate over a specified period of time, and we are interested in the equivalent rate r_β under a different accrual factor β over the same period of time, then this information can be retrieved by making note of the equivalency

$$r_\alpha \alpha = r_\beta \beta. \quad (3.8)$$

This makes conversions simple as demonstrated by the following example.

Example 3.1. If on 16 December 2011 the act/365 rate for three months is 5%, the act/360 rate is

$$r_\beta = r_\alpha \frac{\alpha}{\beta} = 5\% \frac{\text{act}/365}{\text{act}/360} = 5\% \frac{360}{365} = 4.9315\%. \quad (3.9)$$

Example 3.2. The conversion between semi-bonds and annual money is given by

$$\gamma_{SB} \frac{30}{360} = \gamma_{AM} \frac{\text{act}}{360}, \quad (3.10)$$

so that

$$\gamma_{SB} = \gamma_{AM} \frac{\text{act}}{30}, \quad (3.11)$$

and likewise

$$\gamma_{AM} = \gamma_{SB} \frac{30}{\text{act}}. \quad (3.12)$$

Definition 3.9. A *stock* or *share* is an asset giving ownership in a fraction of a company. The price at time t of a stock is denoted by S_t . The current known price is called the *spot*.

Definition 3.10. A *fixed rate bond* with coupon c and notional N is an asset that pays a coupon cN at a specified frequency (usually a year) and N at its maturity date T . A *floating rate bond* is defined similarly with a variable interest rate.

3.1.2 Forward Contracts and Forward Prices

Definition 3.11. A *derivative contract (derivative)* is a financial contract between two parties whose value derives from the value of another variable.

Remark. Derivatives are often defined by the payout at a specified time. E.g. if S_T indicates the total snowfall in inches at time T then the weather derivative $g(S_T)$ can be defined by

$$g(S_T) = I\{S_T > 50\} = \begin{cases} \$1 & \text{if } S_T > 50 \\ 0 & \text{otherwise,} \end{cases} \quad (3.13)$$

where both S_T and $g(S_T)$ are random variables with unknown values until $T = 1$.

Definition 3.12. A *forward contract (forward)* is an agreement between two counterparties to trade a specific asset at maturity T with delivery price K . At a time $t \leq T$, the buyer is *long* the contract while the seller is *short* the contract.

The value at time $t \leq T$ of being long a forward contract is given by $V_K(t, T)$, where by construction we have that $V_K(T, T) = S_T - K$.

Definition 3.13. The *forward price* $F(t, T)$ at current time $t \leq T$ is the delivery price K s.t. $V_K(t, T) = 0$.

Remark. The forward price can be seen as denoting the spot price at maturity, so in essence you will be agreeing at time $t \leq T$ to trade the asset at time T for its actual price. If you had $K < F(t, T)$ then the person long the contract has a better deal since they would pay less than the appreciated value (i.e. you could just sell the asset for profit), whereas $K > F(t, T)$ provides the better deal for the short since they would receive more than the appreciated value.

Table 3.2: Forward prices for various assets

Underlying	Forward price
Asset paying no income	$S_t e^{r(T-t)}$
Asset paying known income I	$(S_t - I) e^{r(T-t)}$
Asset paying dividends at rate q	$S_t e^{(r-q)(T-t)}$
Foreign exchange	$X_t e^{(r_s - r_f)(T-t)}$

The forward price is decided under various conditions, those being listed in Table 3.2. Given the forward price, we derive the value of being long a forward contract as

$$V_K(t, T) = (F(t, T) - K) e^{-r(T-t)}. \quad (3.14)$$

As a final note, this chapter begins to discuss proof methods for deriving the values of derivatives. It's fairly instructional to review these as they form the basis for evaluating many assets, so we shall include the definitions and some proof examples below.

Definition 3.14. A *replication proof* is one where we prove the value of an asset by demonstrating that if two portfolios always have the same value at time T and if neither add/subtract anything of non-zero value between t and T , then they must have the same value at $t \leq T$.

Definition 3.15. *Arbitrage* is a situation where, starting with an empty portfolio, a sequence of simple market transactions will end up with a positive portfolio value at time T .

Definition 3.16. A *no-arbitrage proof* uses the assumption that no arbitrage situations exist, i.e. all market transactions have an equivalent exchange.

Remark. Most no-arbitrage proofs are constructed similarly to a proof by contradiction.

It's often informative constructing no-arbitrage proofs since they force you to think about situations in which you can exploit an arbitrage if detected. We demonstrate no-arbitrage proofs and replication proofs in the following examples.

Example 3.3. Let S_t be the current stock of a price paying no income. Let r_{BID} be the interest rate at which one can invest/lend money, and r_{OFF} be the interest rate at which one can borrow money, $r_{BID} \leq r_{OFF}$. Both rates are continuously compounded. Using arbitrage arguments, find upper and lower bounds for the forward price of the stock for a forward contract with maturity $T > t$.

Solution. Assume that $F(t, T) > S_t e^{r_{OFF}(T-t)}$. At time t , we could short a forward contract with forward price $F(t, T)$ (which is done at no upfront cost since $V_{F(t, T)}(t, T) = 0$), take out a loan for S_t at rate r_{OFF} for time $T-t$, and purchase the asset for S_t . Then at time T , we would execute the forward contract, selling the stock for $F(t, T)$ and paying back the loan with appreciated value $S_t e^{r_{OFF}(T-t)}$, leaving us with $F(t, T) - S_t e^{r_{OFF}(T-t)} > 0$. This provides an upper bound under the assumption of no-arbitrage.

Assume that $F(t, T) < S_t e^{r_{BID}(T-t)}$. At time t , we could long a forward contract with forward price $F(t, T)$, borrow the asset S_t , sell the asset and invest/lend the remains at the rate r_{BID} for time $T-t$. Then at time T , we receive $S_t e^{r_{BID}(T-t)}$ at the maturity of the deposit, which we can use $F(t, T)$ of to execute the forward contract and return the underlying asset. This leaves us with $S_t e^{r_{BID}(T-t)} - F(t, T) > 0$. This provides a lower bound under the assumption of no-arbitrage.

Our bounds then come out to

$$S_t e^{r_{BID}(T-t)} \leq F(t, T) \leq S_t e^{r_{OFF}(T-t)}. \quad (3.15)$$

3.2 A Random Walk Down Wall Street [18]

Part III

Subjects

Chapter 4

Probability and Statistics

Definition 4.1 (Unbiased Estimate). An estimator \hat{y} is called an *unbiased estimator* of y if

$$\mathbb{E} [\hat{y}] = y. \quad (4.1)$$

Chapter 5

Stochastic Optimization

Most stochastic optimization [27] problems follow the form

$$\min_{x \in \mathbb{R}^n} f(x) = \mathbb{E}_{\xi \sim \Xi} [F(x, \xi)] \quad (5.1)$$

where $F : \mathbb{R}^n \times \mathbb{R}^d \rightarrow \mathbb{R}$ is continuously differentiable. Since $F(\cdot, \xi)$ tends to be unknown, this is empirically restated as

$$\min_{x \in \mathbb{R}^n} f(x) = \frac{1}{N} \sum_{i=1}^N f_i(x) \quad (5.2)$$

where $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ is the loss corresponding to the i th sample data, and N denotes the number of sample data and is assumed to be extremely large. We define the following, as they are frequently used as assumptions:

Definition 5.1 (Lipschitz Continuous). A continuously differentiable function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is *Globally Lipschitz Continuous* with constant L if for any $x, y \in \mathbb{R}^n$,

$$\|\nabla f(x) - \nabla f(y)\| \leq L \|x - y\|. \quad (5.3)$$

TODO: Fill out some intuition for Lipschitz Continuity here. Definitely use a figure of sorts.

Remark. In machine learning, we can view x in the previous definitions as being our set of parameters, f as the objective function, and f_i as the component function for data point i (i.e. i corresponds to the i th element in the dataset).

5.1 Quasi-Newton Methods

These methods generally follow the form of Algorithm 4

5.1.1 BFGS

Let $f_k \stackrel{\text{def}}{=} f(\mathbf{x}_k)$. [20] Outlines the baseline method, BFGS, as follows:

Algorithm 4 Stochastic Quasi-Newton Method

Require: $x_1 \in \mathbb{R}^n$, a PD matrix $H_1 \in \mathbb{R}^{n \times n}$, batch sizes $\{m_k\}_{k \geq 1}$, and stepsizes $\{\alpha_k\}_{k \geq 1}$

for $k = 1, 2, \dots$ **do**

- Calculate $g_k = \frac{1}{m_k} \sum_{i=1}^{m_k} g(x_k, \xi_{k,i})$.
- Generate a PD Hessian inverse approximation H_k .
- $x_{k+1} = x_k - \alpha_k H_k g_k$.

end for

Definition 5.2 (Second Order Taylor Polynomial). For a multivariate function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the *Second Order Taylor Polynomial* is given by

$$f(\mathbf{x}) \approx f(\mathbf{x}^*) + \nabla f(\mathbf{x}^*)^\top (\mathbf{x} - \mathbf{x}^*) + \frac{1}{2} (\mathbf{x} - \mathbf{x}^*)^\top H_f(\mathbf{x}^*) (\mathbf{x} - \mathbf{x}^*). \quad (5.4)$$

Definition 5.3 (Wolfe Conditions). A step size α_k is said to satisfy the *Wolfe Conditions* if for a multivariate function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ with iterate $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$, we have

1. $f_{k+1} \leq f_k + c_1 \alpha_k \mathbf{p}_k^\top \nabla f_k$, and
2. $-\mathbf{p}_k^\top \nabla f_{k+1} \leq -c_2 \mathbf{p}_k^\top \nabla f_k$,

where $0 < c_1 < c_2 < 1$.

This section begins by building off of the second order Taylor Polynomial for approximating a function's roots. I will now derive the primary formula used in this paper for my own benefit. Consider the quadratic model at iterate \mathbf{x}_k

$$f(\mathbf{x}) \approx f_k + \nabla f_k^\top (\mathbf{x} - \mathbf{x}_k) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_k)^\top H_{f_k}(\mathbf{x} - \mathbf{x}_k). \quad (5.5)$$

We simplify the notation further by letting $\mathbf{p} = \mathbf{x} - \mathbf{x}_k$ and defining the root finding problem using an approximation of the hessian $\tilde{\mathbf{H}}_k$ at iterate \mathbf{x}_k as

$$m_k(\mathbf{p}) = f_k + \nabla f_k^\top \mathbf{p} + \frac{1}{2} \mathbf{p}^\top \tilde{\mathbf{H}}_k \mathbf{p}. \quad (5.6)$$

The approximate hessian is an $n \times n$ symmetric positive definite matrix that is updated at each iteration.

The minimizer of Equation 5.6 is notably $p_k = -\tilde{\mathbf{H}}_k^{-1} \nabla f_k$, which provides us with the search direction, for which we update the iterate using

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k \quad (5.7)$$

where α_k satisfies the Wolfe Conditions. If f is *strongly convex*, then the *curvature condition*, given by

$$(\mathbf{x}_{k+1} - \mathbf{x}_k)^\top (\nabla f_{k+1} - \nabla f_k) > 0 \quad (5.8)$$

will always hold. Otherwise, we have to force this condition via choice of α_k (directly affects Equation 5.7) since it is a requirement that $\tilde{\mathbf{H}}_{k+1}(\mathbf{x}_{k+1} - \mathbf{x}_k)^\top =$

$\nabla f_{k+1} - \nabla f_k$ in order for the first condition to be met. α_k can be chosen using a line search procedure which imposes the Wolfe or Strong Wolfe conditions.

The paper continues to show that, rather than update the approximate hessian at each iterate, we need only update the inverse since each iterate is a function of the inverse, and each $\nabla m_k(\mathbf{0}) = \nabla f_k$ can be evaluated without the hessian information. Given this, we define the following

$$\mathbf{s}_k = \mathbf{x}_{k+1} - \mathbf{x}_k = \alpha_k \mathbf{p}_k, \quad \mathbf{y}_k = \nabla f_{k+1} - \nabla f_k, \quad \text{and} \quad \rho_k = \frac{1}{\mathbf{y}_k^\top \mathbf{s}_k} \quad (5.9)$$

so that updates to the inverse hessian are given by

$$\tilde{\mathbf{H}}_{k+1}^{-1} = (\mathbf{I} - \rho_k \mathbf{s}_k \mathbf{y}_k^\top) \tilde{\mathbf{H}}_k^{-1} (\mathbf{I} - \rho_k \mathbf{y}_k \mathbf{s}_k^\top) + \rho_k \mathbf{s}_k \mathbf{s}_k^\top. \quad (5.10)$$

5.1.1.1 L-BFGS

TODO: Read LBFGS and write analysis

[27] assumes that f is globally Lipschitz continuous, lower bounded, and that for each iteration k we have an unbiased estimate for the gradient $\mathbb{E}[g(x_k, \xi_k)] = \nabla f(x_k)$. Furthermore, the noise must also be upper-bounded, i.e.

$$\mathbb{E}_{y_k} \left[\|g(x_k, \xi_k) - \nabla f(x_k)\|^2 \right] \leq \sigma^2. \quad (5.11)$$

Additional constraint is placed on the objective function: $F(x, \xi)$ must be twice continuously differentiable w.r.t. x , the stochastic gradient must be computed as $g(x, \xi) = \nabla_x F(x, \xi)$, and there must exist a positive constant κ s.t. $\|\nabla_{xx}^2 F(x, \xi)\| \leq \kappa$ for any x, ξ . This is effectively the same update method as Equation 5.10 (along with L-BFGS which I have yet to write about which approximates $\mathbf{H}_k^{-1} \mathbf{g}_k$ at each iteration rather than solely \mathbf{H}_k^{-1}), except we use a damped \mathbf{y} .

While this method is computationally effective, it is not feasible for storage purposes in deep learning algorithms (as is also the case for L-BFGS). The algorithm requires maintaining all iterates of \mathbf{s}_k and \mathbf{p}_k which is $\mathcal{O}(n)$ space. For purposes such as machine learning, n can be in the millions. This is laid out in lines 5, 6, 10, and 11 of Procedure 3.1 in their paper. Even moreso, the article only compares itself to SGD rather than the more comparable L-BFGS, and so it is hard to tell if its performance even will outdo the method it is built off of.

[6] improves on the L-BFGS method by only maintaining the last m curvature pairs $(\mathbf{y}_k, \mathbf{s}_k)$. Results show that this method performs similarly to that of the fully L-BFGS method even with an overlap of 25%, yet performs significantly more gradient evaluations than that of Adam.

Chapter 6

Federated Learning

[14] outlines the inaugural paper for *federated learning*, a distributed learning method which operates on the basis that most the devices involved in the learning process are heterogenous and have unreliable internet connectivity. The synchronized federated learning algorithm generally follows these steps:

1. a subset of existing clients is selected, each of which downloads the current model;
2. each client in the subset computes an updated model based on their local data;
3. the model updates are sent from the selected clients to the server; and finally
4. the server aggregates these models (typically by averaging) to construct an improved global model.

6.1 Update Methods

The main contributions from [14] deal with deriving a methods which don't require full model updates being sent back and forth between the client and server, for which they derive two primary approaches:

Structured updates Updates are directly learned from a restricted space that can be parameterized using a smaller number of variables.

Sketched updates A full model is learned and then compressed prior to sending it to the server.

The general scheme is to learn a real matrix $\mathbf{W} \in \mathbb{R}^{d_1 \times d_2}$ from data stored across a large number of clients. This is done by sending a current model \mathbf{W}_t to a subset S_t of n_t clients. Each client updates their local model \mathbf{W}_t^i (usually via SGD or some learning algorithm) so that its update is given by $\mathbf{H}_t^i \triangleq \mathbf{W}_i^t - \mathbf{W}_t$. The next global update is then given by

$$\mathbf{W}_{t+1} = \mathbf{W}_t + \eta_t \mathbf{H}_t, \quad \mathbf{H}_t = \frac{1}{n_t} \sum_{i \in S_t} \mathbf{H}_t^i, \quad (6.1)$$

where η_t is the learning rate.

6.1.1 Structured Updates

Structured updates send back and forth full-sized models based on a predefined structure. By doing so, there is little loss of information over the communication channel but potential for higher data throughput.

6.1.1.1 Low Rank

We enforce each $\mathbf{H}_t^i \in \mathbb{R}^{d_1 \times d_2}$ to be a low rank matrix of rank at most k for a fixed k by letting $\mathbf{H}_t^i = \mathbf{A}_t^i \mathbf{B}_t^i$ where $\mathbf{A}_t^i \in \mathbb{R}^{d_1 \times k}$ and $\mathbf{B}_t^i \in \mathbb{R}^{k \times d_2}$. This matrix is further compressed by representing \mathbf{A}_t^i as a random matrix defined by a seed and only optimizing \mathbf{B}_t^i , saving a factor of d_1/k parameters in communication. \mathbf{A}_t^i is generated afresh each round and for each client independently.

6.1.1.2 Random Mask

We restrict each \mathbf{H}_t^i to be sparse based on a predefined sparsity pattern with respect to a random seed (similar to low rank). The sparsity pattern is generated afresh each round and for each client independently.

6.1.2 Sketched Updates

Sketched updates place higher emphasis on communication cost rather than fully structured models.

6.1.2.1 Subsampling

We only communicate a matrix $\hat{\mathbf{H}}_t^i = \mathbf{H}_t^i \odot \mathbf{M}$ where \mathbf{M} is a sparse mask and is generated afresh each round and for each client independently. This can be done s.t. $\mathbb{E} [\hat{\mathbf{H}}_t] = \mathbf{H}_t$.

6.1.2.2 Probabilistic Quantization

Let $\mathbf{h} = (h_1, \dots, h_{d_1 \times d_2}) = \text{vec}(\mathbf{H}_t^i)$, $h_{\max} = \max_j h_j$, and $h_{\min} = \min_j h_j$. We can apply b -bit quantization by dividing $[h_{\min}, h_{\max}]$ into 2^b intervals. Then if h_i falls in the interval $[h_{\min}^i, h_{\max}^i]$, then the compressed update for index i is defined by

$$\tilde{h}_i = \begin{cases} h_{\max}^i, & \text{with probability } \frac{h_i - h_{\min}^i}{h_{\max}^i - h_{\min}^i} \\ h_{\min}^i, & \text{with probability } \frac{h_{\max}^i - h_i}{h_{\max}^i - h_{\min}^i}. \end{cases} \quad (6.2)$$

Using this, we have that $\mathbb{E} [\tilde{\mathbf{h}}] = \mathbf{h}$.

Error bounds on this method can be improved by a factor of $\mathcal{O}(d/\log d)$ by first performing a random rotation on \mathbf{h} , quantizing, sending the information to the server, then performing the inverse rotation. The paper states that a structured rotation matrix, the product of a Walsh-Hadamard and binary diagonal entry

matrix, can have a computational complexity of $\mathcal{O}(d)$ for generating the matrix and $\mathcal{O}(d \log d)$ for applying the matrix.

Chapter 7

Differential Privacy

7.1 (ϵ, δ) -Differential Privacy

Definition 7.1. [7] A randomized algorithm \mathcal{M} is (ϵ, δ) -differentially private if for all $\mathcal{S} \subseteq \text{range}(\mathcal{M})$ and any two neighboring databases with a single modified record $D \sim D'$ has that

$$\Pr [\mathcal{M}(D) \in \mathcal{S}] \leq e^\epsilon \Pr [\mathcal{M}(D') \in \mathcal{S}] + \delta. \quad (7.1)$$

Definition 7.2. [7] The L_n sensitivity of a function f (denoted $\Delta_n(f)$), over any two neighboring databases D, D' (denoted $D \sim D'$) is

$$\Delta_n(f) = \sup_{D \sim D'} \|f(D) - f(D')\|_n. \quad (7.2)$$

7.1.1 Mechanisms

Definition 7.3. The *Randomized Response Mechanism* is defined for a predicate $f : \mathcal{D} \rightarrow \{0, 1\}$ as

$$RR_p(f(D)) \stackrel{\text{def}}{=} \begin{cases} f(D) & \text{with probability } p, \\ 1 - f(D) & \text{with probability } 1 - p. \end{cases} \quad (7.3)$$

Definition 7.4. [7] The *Laplace Mechanism* with parameter ϵ is given by

$$\mathcal{L}_\epsilon(f(\mathbf{X})) \stackrel{\text{def}}{=} f(\mathbf{X}) + (Y_1, \dots, Y_n), \quad (7.4)$$

where each $Y_i \sim \Lambda(0, \Delta_1(f)/\epsilon)$ samples from the *Laplace Distribution* which is characterized by the PDF

$$\Lambda(x|\mu, \sigma) = \frac{1}{2\sigma} \exp\left(-\frac{|x - \mu|}{\sigma}\right). \quad (7.5)$$

Definition 7.5. [7] The *Gaussian Mechanism* with parameter σ is given by

$$\mathcal{G}_\sigma(f(\mathbf{X})) \stackrel{\text{def}}{=} f(\mathbf{X}) + \mathcal{N}(0, \sigma^2 \mathbf{I}). \quad (7.6)$$

7.2 Rényi Differential Privacy

We can improve on the analysis given by (ϵ, δ) -DP using a relaxation provided by Rényi Differential Privacy (RDP). This analysis provides much stronger bounds for our analysis than the strong composition of traditional DP.

Definition 7.6 (Rényi divergence). For two probability distributions P and Q defined over \mathcal{R} , the *Rényi divergence* of order $\alpha > 1$ is

$$D_\alpha(P \parallel Q) \stackrel{\text{def}}{=} \frac{1}{\alpha - 1} \ln \mathbb{E}_{x \in Q} \left[\left(\frac{P(x)}{Q(x)} \right)^\alpha \right]. \quad (7.7)$$

Definition 7.7 $((\alpha, \epsilon)\text{-RDP})$. A randomized mechanism $\mathcal{M} : \mathcal{D} \rightarrow \mathcal{R}$ is said to have ϵ -Rényi differential privacy of order α , or $(\alpha, \epsilon)\text{-RDP}$ for short, if for any adjacent $D, D' \in \mathcal{D}$ it holds that

$$D_\alpha(\mathcal{M}(D) \parallel \mathcal{M}(D')) \leq \epsilon. \quad (7.8)$$

Proposition 7.1 (Composition of RDP). Let $f : \mathcal{D} \rightarrow \mathcal{R}_1$ be $(\alpha, \epsilon_1)\text{-RDP}$ and $g : \mathcal{R}_1 \times \mathcal{D} \rightarrow \mathcal{R}_2$ be $(\alpha, \epsilon_2)\text{-RDP}$. Then the mechanism defined as $(\mathcal{M}_1, \mathcal{M}_2)$, where $\mathcal{M}_1 \sim f(\mathcal{D})$ and $\mathcal{M}_2 \sim g(\mathcal{M}_1, \mathcal{D})$, satisfies $(\alpha, \epsilon_1 + \epsilon_2)\text{-RDP}$.

Definition 7.8 (c -stable Transformations). We say that $g : \mathcal{D} \rightarrow \mathcal{D}'$ is c -stable if $g(A)$ and $g(B)$ are adjacent in \mathcal{D}' implies that there exists a sequence of length $c + 1$ so that $D_0 = A, \dots, D_c = B$ and all (D_i, D_{i+1}) are adjacent in \mathcal{D} .

Proposition 7.2 (2^c -stable RDP). If $f : \mathcal{D} \rightarrow \mathcal{R}$ is $(\alpha, \epsilon)\text{-RDP}$, $g : \mathcal{D}' \rightarrow \mathcal{D}$ is 2^c -stable and $\alpha \geq 2^{c+1}$, then $f \circ g$ is $(\alpha/2^c, 3^c\epsilon)\text{-RDP}$.

Randomized Response RDP

Proposition 7.3. The randomized response mechanism with probability p satisfies

$$\left(\alpha, \frac{1}{\alpha - 1} \log \left(p^\alpha (1-p)^{1-\alpha} + (1-p)^\alpha p^{1-\alpha} \right) \right) - \text{RDP} \quad (7.9)$$

if $\alpha > 1$, and

$$\left(\alpha, (2p - 1) \log \left(\frac{p}{1-p} \right) \right) - \text{RDP} \quad (7.10)$$

if $a = 1$.

Laplacian RDP

Proposition 7.4. For any $\alpha \geq 1$ and $\lambda > 0$:

$$D_\alpha(\Lambda(0, \lambda) \parallel \Lambda(1, \lambda)) = \frac{1}{\alpha - 1} \log \left\{ \frac{\alpha}{2\alpha - 1} \exp \left(\frac{\alpha - 1}{\lambda} \right) + \frac{\alpha - 1}{2\alpha - 1} \exp \left(\frac{-\alpha}{\lambda} \right) \right\}. \quad (7.11)$$

Corollary 7.1. If real-valued function f has sensitivity 1, then the Laplace Mechanism $\mathcal{L}_\lambda(f)$ satisfies $\left(\alpha, \frac{1}{\alpha - 1} \log \left\{ \frac{\alpha}{2\alpha - 1} \exp \left(\frac{\alpha - 1}{\lambda} \right) + \frac{\alpha - 1}{2\alpha - 1} \exp \left(\frac{-\alpha}{\lambda} \right) \right\} \right)$ -RDP.

Gaussian RDP

Proposition 7.5. $D_\alpha(\mathcal{N}(0, \sigma^2) \parallel \mathcal{N}(\mu, \sigma^2)) = \alpha\mu^2/2\sigma^2$

Corollary 7.2. *The gaussian mechanism \mathcal{G}_σ applied to a function f satisfies $(\alpha, \alpha\Delta_2(f)^2/(2\sigma^2))$ -RDP.*

Finally, we can recover the ϵ and δ terms from RDP.

Proposition 7.6. *If \mathcal{M} is an (α, ϵ) -RDP mechanism, it also satisfies $(\epsilon + \frac{\ln 1/\delta}{\alpha-1}, \delta)$ -DP for any $0 < \delta < 1$.*

Chapter 8

Jacobian Orthogonalization in Neural Networks

8.1 Linear Layers

We note that for a fully-connected linear layer $\mathbf{z} = \mathbf{W}\mathbf{h}$, we just need to require and preserve the orthogonality of the weight matrix, since the Jacobian is $\partial\mathbf{z}/\partial\mathbf{h} = \mathbf{W}$. The following methods have been employed for this.

8.1.1 Householder Reflectors

[30] define the SVD of a transition matrix $\mathbf{W} \in \mathbb{R}^{n \times n}$ as given by $\mathbf{W} = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^\top$ where $\boldsymbol{\Sigma}$ is the diagonal matrix of singular values and $\mathbf{U}, \mathbf{V} \in \mathbb{R}^{n \times n}$ are orthogonal.

Definition 8.1. Given a vector $\mathbf{u} \in \mathbb{R}^k$, $k \leq n$, the Householder Reflector $\mathcal{H}_k^n(\mathbf{u})$ is

$$\mathcal{H}_k^n(\mathbf{u}) = \begin{cases} \begin{bmatrix} \mathbf{I}_{n-k} & \mathbf{0} \\ \mathbf{0} & \mathbf{I}_k - 2\frac{\mathbf{u}\mathbf{u}^\top}{\|\mathbf{u}\|^2} \end{bmatrix}, & \mathbf{u} \neq \mathbf{0} \\ \mathbf{I}_n, & \text{otherwise.} \end{cases} \quad (8.1)$$

When n is well-defined, the previous function is shorthanded to $\mathcal{H}_k(\mathbf{u})$.

Let $\mathbb{O}(n)$ be the set of $n \times n$ orthogonal matrices. The following maps are established as precursors to the primary theorems:

$$\begin{aligned} \mathcal{M}_k : \mathbb{R}^k \times \cdots \times \mathbb{R}^n &\rightarrow \mathbb{R}^{n \times n} \\ \mathcal{M}_k(\mathbf{u}_k, \dots, \mathbf{u}_n) &\mapsto \mathcal{H}_n(\mathbf{u}_n) \cdots \mathcal{H}_k(\mathbf{u}_k), \end{aligned} \quad (8.2)$$

and

$$\begin{aligned} \mathcal{M}_{k_1, k_2} : \mathbb{R}^{k_1} \times \cdots \times \mathbb{R}^n \times \mathbb{R}^{k_2} \times \cdots \times \mathbb{R}^n \times \mathbb{R}^n &\rightarrow \mathbb{R}^{n \times n} \\ \mathcal{M}_{k_1, k_2}(\mathbf{u}_{k_1}, \dots, \mathbf{u}_n, \mathbf{v}_{k_1}, \dots, \mathbf{v}_n, \boldsymbol{\sigma}) &\mapsto \\ \mathcal{H}_n(\mathbf{u}_n) \cdots \mathcal{H}_{k_1}(\mathbf{u}_{k_1}) \text{diag}(\boldsymbol{\sigma}) \mathcal{H}_{k_2}(\mathbf{v}_{k_2}) \cdots \mathcal{H}_n(\mathbf{v}_n). \end{aligned} \quad (8.3)$$

Theorem 8.1. $\text{im}(\mathcal{M}_1)$ is the set of all $n \times n$ orthogonal matrices.

Theorem 8.2. $\text{im}(\mathcal{M}_{1,1})$ is the set of all $n \times n$ real matrices.

Theorem 8.3. $\text{im}(\mathcal{M}_{k_1, k_2})$ includes the set of all orthogonal $n \times n$ matrices if $k_1 + k_2 \leq n + 2$.

These definitions can be extended to the case of $\mathbf{W} \in \mathbb{R}^{m \times n}$ by taking

$$\mathbf{W} = \mathbf{U}(\boldsymbol{\Sigma}|0)(\mathbf{V}_L|\mathbf{V}_R)^\top = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}_L^\top, \quad (8.4)$$

where $\mathbf{U} \in \mathbb{R}^{m \times m}$, $\boldsymbol{\Sigma} \in \text{diag}(\mathbb{R}^m)$, and $\mathbf{V}_L \in \mathbb{R}^{n \times m}$. Equation 8.3 can then be extended as

$$\begin{aligned} \mathcal{M}_{k_1, k_2}^{m, n} : \mathbb{R}^{k_1} \times \cdots \times \mathbb{R}^m \times \mathbb{R}^{k_2} \times \cdots \times \mathbb{R}^n \times \mathbb{R}^{\min(m, n)} &\rightarrow \mathbb{R}^{m \times n} \\ \mathcal{M}_{k_1, k_2}^{m, n}(\mathbf{u}_{k_1}, \dots, \mathbf{u}_m, \mathbf{v}_{k_1}, \dots, \mathbf{v}_n, \sigma) &\mapsto \\ \mathcal{H}_m^m(\mathbf{u}_m) \cdots \mathcal{H}_{k_1}^m(\mathbf{u}_{k_1}) \hat{\boldsymbol{\Sigma}} \mathcal{H}_{k_2}^n(\mathbf{v}_{k_2}) \cdots \mathcal{H}_n^n(\mathbf{v}_n). \end{aligned} \quad (8.5)$$

where $\hat{\boldsymbol{\Sigma}} = (\text{diag}(\boldsymbol{\sigma})|0)$ if $m < n$ and $(\text{diag}(\boldsymbol{\sigma})|0)^\top$ otherwise. This leads to the following lemma and consequently theorem:

Lemma 8.1. Given $\{v_i\}_{i=1}^n$, define $V^{(k)} = \mathcal{H}_n^n(\mathbf{v}_n) \cdots \mathcal{H}_k^n(\mathbf{v}_k)$ for $k \in [n]$. We have:

$$V_{*, i}^{(k_1)} = V_{*, i}^{(k_2)}, \forall k_1, k_2 \in [n], i \leq \min(n - k_1, n - k_2). \quad (8.6)$$

Theorem 8.4. If $m \leq n$, $\text{im}(\mathcal{M}_{1, n-m+1}^{m, n})$ is the set of all $m \times n$ matrices; else the image of $\text{im}(\mathcal{M}_{m-n+1, 1}^{m, n})$ is the set of all $m \times n$ matrices.

8.1.2 Orthogonality Regularization

[28] uses an orthogonality regularizer, which is done by maintaining each weight matrix's SVD $\mathbf{W} = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^\top$, and enforcing a regularizer as part of the objective function, given by

$$L_o(\mathbf{U}, \mathbf{V}) = \frac{1}{r^2} \left(\|\mathbf{U}^\top \mathbf{U} - \mathbf{I}\|_F^2 + \|\mathbf{V}^\top \mathbf{V} - \mathbf{I}\|_F^2 \right). \quad (8.7)$$

Using this method, \mathbf{W} will maintain an approximate orthogonality so long as we have all singular values $\sigma_i(\boldsymbol{\Sigma}) \in \{-1, 1\}$

8.1.3 Cayley Transform

The *Cayley Transform*, used in [26, 29], make use of skew symmetric matrices, i.e. matrices s.t. $\mathbf{A} = -\mathbf{A}^\top$, to construct their orthogonal convolutions since a skew symmetric matrix can be constructed for any matrix \mathbf{B} using the formula $\mathbf{A} = \mathbf{B} - \mathbf{B}^\top$. From this, an orthogonal matrix is explicitly constructed via $\mathbf{Q} = (\mathbf{I} - \mathbf{A})(\mathbf{I} + \mathbf{A})^{-1}$.

8.2 Convolutional Layers

In this section, we consider the case of convolutional filters having an orthogonal Jacobian. A simple way of viewing this comes from that fact that for any convolutional filter \mathbf{C} , the convolution operation $\mathbf{Y} = \mathbf{C} \star \mathbf{X}$ is a linear transformation, and thus there exists some $\mathbf{W}_\mathbf{C}$ such that

$$\mathbf{Y} = \mathbf{C} \star \mathbf{X} \Leftrightarrow \vec{\mathbf{Y}} = \mathbf{W}_\mathbf{C} \vec{\mathbf{X}}. \quad (8.8)$$

We then say that \mathbf{C} is orthogonal if there exists an orthogonal $\mathbf{W}_\mathbf{C}$.

8.2.1 Explicit Convolution Construction

Using the fact that \mathbf{C} is equivalent to some $\mathbf{W}_\mathbf{C}$, we can explicitly construct $\mathbf{W}_\mathbf{C}$ as orthogonal so that by definition, \mathbf{C} has an orthogonal Jacobian. We may then perform the transformation via matrix-vector multiplication on $\vec{\mathbf{X}}$ and reshape the resultant vector $\vec{\mathbf{Y}}$ to be that of the desired output. Using this, the methods employed in Section 8.1 are viable; however, they are often computationally infeasible due to the size of convolution filters.

8.2.2 Circularly Padded Convolutions

The circular convolution \mathbf{C} has the notable property that its Jacobian is circulant, and that the matrix form of this convolution $\mathbf{W}_\mathbf{C}$ is a circulant matrix. [17, 26, 23] all require circularly padded convolutions, as [23, Appendix D.2] has proved the alternative zero-padded orthogonal convolutions must be of size 1.

[26] lay out their framework using the *2D convolution theorem* [12] and the *circular convolution-multiplication property* [13], given in their analysis by

$$(\mathbf{W} \star \mathbf{X})_{:,i,j} = \text{IFFT} \left(\text{FFT}(\mathbf{W})_{:,:,i,j} \text{FFT}(\mathbf{x})_{:,:,i,j} \right). \quad (8.9)$$

Since the convolution in the Fourier domain is just the element-wise product, they can efficiently enforce each $\text{FFT}(\mathbf{W})_{:,:,i,j}$ as being orthogonal via the Cayley transform method highlighted in Subsection 8.1.3. This allows them to perform the calculation much more efficiently than if they explicitly constructed the corresponding weight matrix $\mathbf{W}_\mathbf{C}$. Since inverse convolutions are necessary for this method, natural strided convolutions are incompatible and so invertible down-sampling is necessary to emulate striding. Additionally, the parameterization is unable to obtain orthogonal convolutions with -1 eigenvalues, so the whole space of orthogonal convolutions is not represented.

[23] makes use of a modification on skew symmetric matrices (defined in Subsection 8.1.3), which for any convolution filter \mathbf{C} can be constructed by $\mathbf{M} = \mathbf{C} - \text{conv_transpose}(\mathbf{C})$. From their, they build their orthogonal Jacobian off of the *convolutional exponential* [11] which is defined for a skew-symmetric convolutional filter $\mathbf{L} \in \mathbb{R}^{m \times m \times k \times k}$ and input $\mathbf{X} \in \mathbb{R}^{m \times n \times n}$ as

$$\mathbf{L} \star_e \mathbf{X} = \mathbf{X} + \sum_{i=1}^{\infty} \frac{\mathbf{L} \star^i \mathbf{X}}{i!} \quad (8.10)$$

where it's assumed $m = c_{\text{in}} = c_{\text{out}}$. When $c_{\text{in}} \neq c_{\text{out}}$, other cases are considered. By construction, the Jacobian of the loss with respect to this operation satisfies $\exp(\mathbf{J}) \vec{\mathbf{X}} = \overrightarrow{\mathbf{L} *_e \mathbf{X}}$ where \mathbf{L} is skew symmetric which they use to show that \mathbf{J} (the Jacobian of $\mathbf{L} * \mathbf{X}$) is skew symmetric, and therefore $\exp(\mathbf{J})$ is orthogonal. The paper approximates $\exp(\mathbf{J})$ by taking the first k terms, denoted $\mathbf{S}_k(\mathbf{J})$ which yields an approximation with error $\|\exp(\mathbf{J}) - \mathbf{S}_k(\mathbf{J})\| \leq \|\mathbf{J}\|_2^k / k!$ [23, Theorem 3].

While this provides a SOA method for explicitly constructing orthogonal matrices from Skew Symmetric ones, and outperforms the methods defined in BCOP [17], its evaluation time suffers due to the repeated application of the convolution filter to approximate the orthogonal matrix. However, this method forms the basis for that of [29], which either uses either [26] or [23] to construct each kernel as orthogonal , and conducting a dilated convolution using the recovered kernels to get the orthogonal convolution.

Chapter 9

Random Projections

Definition 9.1. *Random projections* are low-dimensional embeddings $\Pi : \mathbb{R}^p \rightarrow \mathbb{R}^\tau$ which preserve - up to a small distortion - the geometry of a subspace of vectors.

Definition 9.2. The *Subsampled Randomized Fourier Transform* is a *structured* random projection given by

$$\Pi = \sqrt{p/\tau} \mathbf{S} \boldsymbol{\Theta} \mathbf{D} \quad (9.1)$$

where

1. $\mathbf{S} \in \mathbb{R}^{\tau \times p}$ is a subsampling matrix,
2. $\boldsymbol{\Theta} \in \mathbb{R}^{p \times p}$ is a unitary discrete fourier transform matrix.
3. $\mathbf{D} \in \mathbb{R}^{p \times p}$ is a diagonal matrix whose entries are drawn independently from $\{-1, 1\}$,

Bibliography

- [1] Gagan Aggarwal et al. “Approximation Algorithms for k-Anonymity”. In: *Proceedings of the International Conference on Database Theory (ICDT 2005)*. Nov. 2005. URL: <http://ilpubs.stanford.edu:8090/645/>.
- [2] Stefan Arnborg, Derek G. Corneil, and Andrzej Proskurowski. “Complexity of Finding Embeddings in a k-Tree”. In: *SIAM Journal on Algebraic Discrete Methods* 8.2 (1987), pp. 277–284. DOI: 10.1137/0608024. eprint: <https://doi.org/10.1137/0608024>. URL: <https://doi.org/10.1137/0608024>.
- [3] Achraf Bahamou, Donald Goldfarb, and Yi Ren. *A Mini-Block Natural Gradient Method for Deep Neural Networks*. 2022. DOI: 10.48550/ARXIV.2202.04124. URL: <https://arxiv.org/abs/2202.04124>.
- [4] S. Blyth. *An Introduction to Quantitative Finance*. OUP Oxford, 2013. ISBN: 9780199666591. URL: <https://books.google.com/books?id=SXbcAAAAQBAJ>.
- [5] Hans L Bodlaender and Rolf H Möhring. “The pathwidth and treewidth of cographs”. In: *SIAM Journal on Discrete Mathematics* 6.2 (1993), pp. 181–188.
- [6] Raghu Bollapragada et al. “A progressive batching L-BFGS method for machine learning”. In: *International Conference on Machine Learning*. PMLR. 2018, pp. 620–629.
- [7] Cynthia Dwork and Aaron Roth. “The Algorithmic Foundations of Differential Privacy”. In: *Found. Trends Theor. Comput. Sci.* 9.3–4 (Aug. 2014), pp. 211–407. ISSN: 1551-305X. DOI: 10.1561/0400000042. URL: <https://doi.org/10.1561/0400000042>.
- [8] Michel X. Goemans and David P. Williamson. “Improved Approximation Algorithms for Maximum Cut and Satisfiability Problems Using Semidefinite Programming”. In: *J. ACM* 42.6 (Nov. 1995), pp. 1115–1145. ISSN: 0004-5411. DOI: 10.1145/227683.227684. URL: <https://doi.org/10.1145/227683.227684>.
- [9] Ian Goodfellow. *Efficient Per-Example Gradient Computations*. 2015.
- [10] Pinar Heggernes. *Treewidth, partial k-trees, and chordal graphs*. Sept. 2006.
- [11] Emiel Hoogeboom et al. “The Convolution Exponential and Generalized Sylvester Flows”. In: *CoRR* abs/2006.01910 (2020). arXiv: 2006.01910. URL: <https://arxiv.org/abs/2006.01910>.

- [12] Anil K Jain. *Fundamentals of digital image processing*. Prentice-Hall, Inc., 1989.
- [13] Mahdi Karami et al. “Invertible convolutional flow”. In: *Advances in Neural Information Processing Systems* 32 (2019).
- [14] Jakub Konečný et al. “Federated Learning: Strategies for Improving Communication Efficiency”. In: *CoRR* abs/1610.05492 (2016). arXiv: 1610.05492. URL: <http://arxiv.org/abs/1610.05492>.
- [15] Gabriel Krummenacher et al. *Scalable Adaptive Stochastic Optimization Using Random Projections*. 2016. arXiv: 1611.06652 [stat.ML].
- [16] Jan van Leeuwen. “Graph algorithms”. In: *Algorithms and complexity*. Elsevier, 1990, pp. 525–631.
- [17] Qiyang Li et al. “Preventing Gradient Attenuation in Lipschitz Constrained Convolutional Networks”. In: *Advances in Neural Information Processing Systems*. Ed. by H. Wallach et al. Vol. 32. Curran Associates, Inc., 2019. URL: <https://proceedings.neurips.cc/paper/2019/file/1ce3e6e3f452828e23a0c94572bef9Paper.pdf>.
- [18] Burton. G. Malkiel. *A Random Walk Down Wall Street*. Norton, New York, 1973.
- [19] James Martens. “New insights and perspectives on the natural gradient method”. In: *arXiv preprint arXiv:1412.1193* (2014).
- [20] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. 2e. New York, NY, USA: Springer, 2006.
- [21] Jiang Qian et al. “Understanding gradient clipping in incremental gradient methods”. In: *International Conference on Artificial Intelligence and Statistics*. PMLR. 2021, pp. 1504–1512.
- [22] Donald J. Rose. “On Simple Characterizations of K-Trees”. In: *Discrete Math.* 7.3–4 (Jan. 1974), pp. 317–322. ISSN: 0012-365X. DOI: 10.1016/0012-365X(74)90042-9. URL: [https://doi.org/10.1016/0012-365X\(74\)90042-9](https://doi.org/10.1016/0012-365X(74)90042-9).
- [23] Sahil Singla and Soheil Feizi. “Skew Orthogonal Convolutions”. In: *CoRR* abs/2105.11417 (2021). arXiv: 2105.11417. URL: <https://arxiv.org/abs/2105.11417>.
- [24] Hidenori Tanaka et al. “Pruning neural networks without any data by iteratively conserving synaptic flow”. In: *CoRR* abs/2006.05467 (2020). arXiv: 2006.05467. URL: <https://arxiv.org/abs/2006.05467>.
- [25] Roberto Tempo and Hideaki Ishii. “Monte Carlo and Las Vegas Randomized Algorithms for Systems and Control*: An Introduction”. In: *European Journal of Control* 13.2 (2007), pp. 189–203. ISSN: 0947-3580. DOI: <https://doi.org/10.3166/ejc.13.189-203>. URL: <https://www.sciencedirect.com/science/article/pii/S0947358007708191>.
- [26] Asher Trockman and J. Zico Kolter. “Orthogonalizing Convolutional Layers with the Cayley Transform”. In: *CoRR* abs/2104.07167 (2021). arXiv: 2104.07167. URL: <https://arxiv.org/abs/2104.07167>.

- [27] Xiao Wang et al. *Stochastic Quasi-Newton Methods for Nonconvex Stochastic Optimization*. 2017. arXiv: 1607.01231 [math.OC].
- [28] Huanrui Yang et al. “Learning Low-rank Deep Neural Networks via Singular Vector Orthogonality Regularization and Singular Value Sparsification”. In: *CoRR* abs/2004.09031 (2020). arXiv: 2004 . 09031. URL: <https://arxiv.org/abs/2004.09031>.
- [29] Tan Yu et al. “Constructing Orthogonal Convolutions in an Explicit Manner”. In: *International Conference on Learning Representations*. 2022. URL: <https://openreview.net/forum?id=Zr5W2LSRhD>.
- [30] Jiong Zhang, Qi Lei, and Inderjit Dhillon. “Stabilizing Gradients for Deep Neural Networks via Efficient SVD Parameterization”. In: *Proceedings of the 35th International Conference on Machine Learning*. Ed. by Jennifer Dy and Andreas Krause. Vol. 80. Proceedings of Machine Learning Research. PMLR, July 2018, pp. 5806–5814. URL: <https://proceedings.mlr.press/v80/zhang18g.html>.
- [31] Xiangyu Zhang et al. “Accelerating Very Deep Convolutional Networks for Classification and Detection”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 38.10 (2016), pp. 1943–1955. DOI: 10.1109/TPAMI.2015.2502579.